

# main

October 14, 2023

## 1 Data Preprocessing Project

### 1.1 Team members

Phuc Dinh  
Brian Shao  
Steven Ho  
Alex Souv  
Navteg Khalsa  
Moshley Marcelo

### 1.2 Import required modules

```
[ ]: from scipy.stats import zscore
     from sklearn.model_selection import train_test_split
     import numpy as np
     import pandas as pd
```

### 1.3 Data Preprocessing

#### 1.3.1 Dataset 1 (London Air)

In this dataset, we use these techniques: dropping fields, removing rows with missing values

#### Import data

```
[ ]: laqndata = pd.read_csv('https://raw.githubusercontent.com/dinhphucv/CSC-177/
    ↪main/Data%20Preprocessing%20Project/data/LaqnData.csv')
```

```
[ ]: laqndata.iloc[0:5]
```

```
[ ]: 
```

	Site	Species	ReadingDateTime	Value	Units	Provisional	or Ratified
0	HI0	CO	01/01/2022 00:00	NaN	mg m-3		P
1	HI0	CO	01/01/2022 00:15	NaN	mg m-3		P
2	HI0	CO	01/01/2022 00:30	NaN	mg m-3		P
3	HI0	CO	01/01/2022 00:45	NaN	mg m-3		P
4	HI0	CO	01/01/2022 01:00	NaN	mg m-3		P

**Remove unnecessary columns** This piece of code allows us to remove column 'Provisional or Ratified' which is not necessary.

```
[ ]: laqndata.drop(columns=['Provisional or Ratified'], axis=1, inplace=True)
```

**Drop rows with missing values** As we can see from the tabular above, some values are NaN. The code below is used to drop all rows with missing values

```
[ ]: print('Number of rows in original data:', laqndata.shape[0])
laqndata.dropna(inplace=True)
laqndata
print('Number of rows after removing NaN:', laqndata.shape[0])
```

Number of rows in original data: 175200  
Number of rows after removing NaN: 104596

```
[ ]: laqndata.iloc[0:5]
```

```
[ ]:      Site Species   ReadingDateTime  Value  Units
35040  HIO        NO  01/01/2022 00:00    2.4  ug m-3
35041  HIO        NO  01/01/2022 00:15    2.4  ug m-3
35042  HIO        NO  01/01/2022 00:30    2.4  ug m-3
35043  HIO        NO  01/01/2022 00:45    2.4  ug m-3
35044  HIO        NO  01/01/2022 01:00    1.2  ug m-3
```

### 1.3.2 Dataset 2 (UNSW-NB15)

In this dataset, we use these techniques to clean the dataset: dropping fields, removing duplicated rows

#### Import data

```
[ ]: unsw_nb15 = pd.read_csv('https://github.com/dinhphucv/CSC-177/raw/main/
↳Data%20Preprocessing%20Project/data/UNSW_NB15.csv')
```

```
[ ]: unsw_nb15.iloc[0:5]
```

```
[ ]:      id      dur proto service state  spkts  dpkts  sbytes  dbytes    rate  \
0    1  0.121478  tcp      -   FIN        6     4    258    172  74.087490
1    2  0.649902  tcp      -   FIN       14    38    734   42014  78.473372
2    3  1.623129  tcp      -   FIN        8    16    364   13186  14.170161
3    4  1.681642  tcp    ftp   FIN       12    12    628     770  13.677108
4    5  0.449454  tcp      -   FIN       10     6    534     268  33.373826

      ...  ct_dst_sport_ltm  ct_dst_src_ltm  is_ftp_login  ct_ftp_cmd  \
0    ...                1                1            0            0
1    ...                1                2            0            0
2    ...                1                3            0            0
3    ...                1                3            1            1
```

	...	1	40	0	0
	ct_flw_http_mthd	ct_src_ltm	ct_srv_dst	is_sm_ips_ports	attack_cat \
0	0	1	1	0	Normal
1	0	1	6	0	Normal
2	0	2	6	0	Normal
3	0	2	1	0	Normal
4	0	2	39	0	Normal

	label
0	0
1	0
2	0
3	0
4	0

[5 rows x 45 columns]

**Check and save 'unsw\_nb15' dataset shape**

```
[ ]: unsw_nb15_original_shape = unsw_nb15.shape
      print('Number of instances/records:', unsw_nb15_original_shape[0])
      print('Number of features/attributes:', unsw_nb15_original_shape[1])
```

Number of instances/records: 175341

Number of features/attributes: 45

**Remove unnecessary columns** The code below allows us to remove multiple categorical ('id' and 'attack\_cat') as we are only interested in whether the connect is an attack rather than the attack type. Also, the id shouldn't have any meaning or effect on the output.

```
[ ]: unsw_nb15.drop(columns=['id', 'attack_cat'], axis=1, inplace=True)
```

**Check and remove duplicated rows** Here, we remove all duplicated rows. We have to do this after removing 'id' columns as the id may prevent us from remove duplicating data.

```
[ ]: print('Duplicated rows in \'unsw_nb15\' dataset:', unsw_nb15.duplicated().sum())
      if unsw_nb15.duplicated().sum() > 0:
          unsw_nb15.drop_duplicates(inplace=True)
```

Duplicated rows in 'unsw\_nb15' dataset: 74072

```
[ ]: print('\''unsw_nb15\' original shape:', unsw_nb15_original_shape)
      print('\''unsw_nb15\' current shape:', unsw_nb15.shape)
```

'unsw\_nb15' original shape: (175341, 45)

'unsw\_nb15' current shape: (101269, 43)

### 1.3.3 Dataset 3 (diabetes)

In this dataset, we use these techniques: shuffling, sorting, saving dataframe

#### Import data

```
[ ]: diabetes = pd.read_csv('https://raw.githubusercontent.com/dinhphucv/CSC-177/
    ↪main/Data%20Preprocessing%20Project/data/diabetes.csv')
```

```
[ ]: diabetes.iloc[0:5]
```

```
[ ]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0             6      148            72           35          0  33.6
1             1       85            66           29          0  26.6
2             8      183            64            0          0  23.3
3             1       89            66           23         94  28.1
4             0      137            40           35        168  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
1                0.351    31         0
2                0.672    32         1
3                0.167    21         0
4                2.288    33         1
```

#### Check and save 'diabetes' dataset shape

```
[ ]: diabetes_original_shape = diabetes.shape
print('Number or instances/records:', diabetes_original_shape[0])
print('Number or features/attributes:', diabetes_original_shape[1])
```

Number or instances/records: 768

Number or features/attributes: 9

#### Check and remove duplicated rows

```
[ ]: print('Duplicated rows in \'diabetes\' dataset:', diabetes.duplicated().sum())
```

Duplicated rows in 'diabetes' dataset: 0

#### Outliers

```
[ ]: diabetes[diabetes.columns.difference(['Outcome'])].info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 768 entries, 0 to 767

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Age	768 non-null	int64
1	BMI	768 non-null	float64
2	BloodPressure	768 non-null	int64
3	DiabetesPedigreeFunction	768 non-null	float64


```

4   Glucose          768 non-null   int64
5   Insulin          768 non-null   int64
6   Pregnancies      768 non-null   int64
7   SkinThickness    768 non-null   int64
dtypes: float64(2), int64(6)
memory usage: 48.1 KB

```

```
[ ]: diabetes_2 = diabetes[diabetes.columns.difference(['Outcome'])].copy()
diabetes_2.boxplot(figsize=(20,10))
```

```
[ ]: <Axes: >
```



The boxplots suggest that there're some abnormally high values in all columns. The code below shows the results of standardizing the columns of the data. We also discarding the outlier with Z score  $> 3$  or  $Z \leq -3$

```
[ ]: Z = (diabetes_2-diabetes_2.mean())/diabetes_2.std()
Z
```

```
[ ]:
      Age      BMI  BloodPressure  DiabetesPedigreeFunction  Glucose  \
0   1.425067  0.203880      0.149543          0.468187  0.847771
1  -0.190548 -0.683976     -0.160441         -0.364823 -1.122665
2  -0.105515 -1.102537     -0.263769          0.604004  1.942458
3  -1.040871 -0.493721     -0.160441         -0.920163 -0.997558
4  -0.020483  1.408828     -1.503707          5.481337  0.503727
..      ...      ...           ...           ...           ...
763  2.530487  0.115094      0.356200         -0.908090 -0.622237
764 -0.530677  0.609757      0.046215         -0.398023  0.034575
765 -0.275580 -0.734711      0.149543         -0.684747  0.003299

```

```

766  1.169970 -0.240048      -0.470426      -0.370859  0.159683
767 -0.870806 -0.201997      0.046215      -0.473476 -0.872451

```

```

      Insulin  Pregnancies  SkinThickness
0   -0.692439    0.639530    0.906679
1   -0.692439   -0.844335    0.530556
2   -0.692439    1.233077   -1.287373
3    0.123221   -0.844335    0.154433
4    0.765337   -1.141108    0.906679
..      ...      ...      ...
763  0.869464    1.826623    1.721613
764 -0.692439   -0.547562    0.405181
765  0.279412    0.342757    0.154433
766 -0.692439   -0.844335   -1.287373
767 -0.692439   -0.844335    0.655930

```

[768 rows x 8 columns]

```

[ ]: Z2 = Z.loc[((Z > -3).sum(axis=1)==8) & ((Z <= 3).sum(axis=1)==8),:]
Z2

```

```

[ ]:      Age      BMI  BloodPressure  DiabetesPedigreeFunction  Glucose  \
0    1.425067  0.203880    0.149543      0.468187  0.847771
1   -0.190548 -0.683976   -0.160441     -0.364823 -1.122665
2   -0.105515 -1.102537   -0.263769      0.604004  1.942458
3   -1.040871 -0.493721   -0.160441     -0.920163 -0.997558
5   -0.275580 -0.810813    0.252871     -0.817546 -0.153085
..      ...      ...      ...      ...      ...
763  2.530487  0.115094    0.356200     -0.908090 -0.622237
764 -0.530677  0.609757    0.046215     -0.398023  0.034575
765 -0.275580 -0.734711    0.149543     -0.684747  0.003299
766  1.169970 -0.240048   -0.470426     -0.370859  0.159683
767 -0.870806 -0.201997    0.046215     -0.473476 -0.872451

```

```

      Insulin  Pregnancies  SkinThickness
0   -0.692439    0.639530    0.906679
1   -0.692439   -0.844335    0.530556
2   -0.692439    1.233077   -1.287373
3    0.123221   -0.844335    0.154433
5   -0.692439    0.342757   -1.287373
..      ...      ...      ...
763  0.869464    1.826623    1.721613
764 -0.692439   -0.547562    0.405181
765  0.279412    0.342757    0.154433
766 -0.692439   -0.844335   -1.287373
767 -0.692439   -0.844335    0.655930

```

[688 rows x 8 columns]

**Shuffling dataframe** Here, we are shuffling dataframe

```
[ ]: # np.random.seed(38)
temp_df = diabetes
temp_df = temp_df.reindex(np.random.permutation(temp_df.index))
temp_df.reset_index(inplace=True, drop=True)
temp_df
```

```
[ ]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0                3    116           74           15      105  26.3
1                5    115           76            0         0  31.2
2                7    168           88           42     321  38.2
3                6    190           92            0         0  35.5
4                8     91           82            0         0  35.6
..            ...    ...           ...           ...    ...  ...
763              8    194           80            0         0  26.1
764              6     80           66           30         0  26.2
765              4     95           60           32         0  35.4
766              6    108           44           20     130  24.0
767              1     89           66           23         94  28.1
```

```
      DiabetesPedigreeFunction  Age  Outcome
0                0.107    24         0
1                0.343    44         1
2                0.787    40         1
3                0.278    66         1
4                0.587    68         0
..            ...    ...    ...
763              0.551    67         0
764              0.313    41         0
765              0.284    28         0
766              0.813    35         0
767              0.167    21         0
```

[768 rows x 9 columns]

**Sorting dataframe** The below code allows us to sort the data by 'Age' in ascending.

```
[ ]: temp_df = diabetes.sort_values(by='Age', ascending=True)
temp_df
```

```
[ ]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
255                1    113           64           35         0  33.6
60                 2     84            0            0         0   0.0
102                0    125           96            0         0  22.5
```

182	1	0	74	20	23	27.7
623	0	94	70	27	115	43.5
..	...	...	...	...	...	...
123	5	132	80	0	0	26.8
684	5	136	82	0	0	0.0
666	4	145	82	18	0	32.5
453	2	119	0	0	0	19.6
459	9	134	74	33	60	25.9

	DiabetesPedigreeFunction	Age	Outcome
255	0.543	21	1
60	0.304	21	0
102	0.262	21	0
182	0.299	21	0
623	0.347	21	0
..	...	...	...
123	0.186	69	0
684	0.640	69	0
666	0.235	70	1
453	0.832	72	0
459	0.460	81	0

[768 rows x 9 columns]

**Saving dataframe for later use** After doing some data preprocessing, we can export the dataframe as csv for later use. We can continue the work instead of redone it.

```
[ ]: temp_df.to_csv('data/diabetes-age-sorted.csv', index=False)
```

### 1.3.4 Dataset 4 (Mobile Phone Price)

In this dataset, we use these techniques: removing duplicated rows, concatenating, calculated fields

#### Import data

```
[ ]: mobile_phone_price = pd.read_csv('https://raw.githubusercontent.com/dinhphucv/
↳CSC-177/main/Data%20Preprocessing%20Project/data/Mobile_Phone_Price.csv')
mobile_phone_price
```

```
[ ]:
      Brand      Model Storage  RAM  Screen Size (inches) \
0    Apple    iPhone 13 Pro  128 GB   6 GB              6.1
1  Samsung  Galaxy S21 Ultra  256 GB  12 GB              6.8
2  OnePlus      9 Pro    128 GB   8 GB              6.7
3  Xiaomi  Redmi Note 10 Pro  128 GB   6 GB             6.67
4   Google      Pixel 6    128 GB   8 GB              6.4
..     ...      ...      ...      ...              ...
402 Samsung  Galaxy Note20 5G   128    8              6.7
403  Xiaomi      Mi 10 Lite 5G   128    6             6.57
```



404	Apple	iPhone 12 Pro Max	128	6	6.7
405	Oppo	Reno3	128	8	6.4
406	Samsung	Galaxy S10 Lite	128	6	6.7

	Camera (MP)	Battery Capacity (mAh)	Price (\$)
0	12 + 12 + 12	3095	999
1	108 + 10 + 10 + 12	5000	1199
2	48 + 50 + 8 + 2	4500	899
3	64 + 8 + 5 + 2	5020	279
4	50 + 12.2	4614	799
..	...	...	...
402	12+64+12	4300	1049
403	48+8+2+2	4160	349
404	12+12+12	3687	1099
405	48+13+8+2	4025	429
406	48+12+5	4500	649

[407 rows x 8 columns]

**Concatenating** In this part, we replicate concatenating by creating multiple dataframes and concatenating them into one.

```
[ ]: col_brand = mobile_phone_price['Brand']
      col_model = mobile_phone_price['Model']
      col_storage = mobile_phone_price['Storage ']
      col_ram = mobile_phone_price['RAM ']
      col_price = mobile_phone_price['Price ($)']
```

```
[ ]: result = pd.concat([col_brand, col_model, col_storage, col_ram,
      ↪col_price],axis=1)
result
```

	Brand	Model	Storage	RAM	Price (\$)
0	Apple	iPhone 13 Pro	128 GB	6 GB	999
1	Samsung	Galaxy S21 Ultra	256 GB	12 GB	1199
2	OnePlus	9 Pro	128 GB	8 GB	899
3	Xiaomi	Redmi Note 10 Pro	128 GB	6 GB	279
4	Google	Pixel 6	128 GB	8 GB	799
..	...	...	...	...	...
402	Samsung	Galaxy Note20 5G	128	8	1049
403	Xiaomi	Mi 10 Lite 5G	128	6	349
404	Apple	iPhone 12 Pro Max	128	6	1099
405	Oppo	Reno3	128	8	429
406	Samsung	Galaxy S10 Lite	128	6	649

[407 rows x 5 columns]

**Removing certain characters in certain columns** As we see from the above tabular, the values for Storage, RAM, and Price are not formatted consistently. Some have ‘GB’, and some do not. The same thing happened with the Price column but with the character ‘\$’. The below code is used to remove that inconsistency.

```
[ ]: result.info()
for column in ['Storage ', 'RAM ']:
    result[column] = result[column].str.replace(r' ?GB', '', regex=True)
result['Price ($)'] = result['Price ($)'].str.replace('[$,]', '', regex=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 407 entries, 0 to 406
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Brand       407 non-null   object
1   Model       407 non-null   object
2   Storage     407 non-null   object
3   RAM         407 non-null   object
4   Price ($)   407 non-null   object
dtypes: object(5)
memory usage: 16.0+ KB
```

**Convert to numeric data type** After removing the inconsistent, the dtype of those columns is still an object. The below code is used to convert them back to numeric. If any non-numerics are still left in those columns, errors will occur. If it ran successfully, we converted them to numeric data type.

```
[ ]: for column in ['Storage ', 'RAM ', 'Price ($)']:
    result[column] = pd.to_numeric(result[column])
result.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 407 entries, 0 to 406
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Brand       407 non-null   object
1   Model       407 non-null   object
2   Storage     407 non-null   int64
3   RAM         407 non-null   int64
4   Price ($)   407 non-null   int64
dtypes: int64(3), object(2)
memory usage: 16.0+ KB
```

**Check and remove duplicated rows** In this step, we remove duplicated rows. We decided to do this after removing inconsistent and converting as there’s a chance the collector accidentally recollected the data and added inconsistent to it.

```
[ ]: print('Duplicated rows in \'result\' dataset:', result.duplicated().sum())
      if result.duplicated().sum() > 0:
          result.drop_duplicates(inplace=True)
```

Duplicated rows in 'result' dataset: 60

**Calculated Fields** After removing all the duplicated, we can add a calculated field into the current dataframe. Here we add a new column 'Price (in VND)' and it will be calculated from 'Price (\$)'. In this piece of code, we also use round() to round the output as Vietnamese currency is always integer.

```
[ ]: result.insert(len(result.columns), 'Price (VND)', (result['Price ($)']*24432.50).
      ↪round().astype(int))
      result
```

```
[ ]:      Brand          Model  Storage  RAM  Price ($)  Price (VND)
0      Apple      iPhone 13 Pro      128    6         999      24408068
1      Samsung  Galaxy S21 Ultra      256   12        1199      29294568
2      OnePlus          9 Pro      128    8         899      21964818
3      Xiaomi  Redmi Note 10 Pro      128    6         279       6816668
4      Google          Pixel 6      128    8         799      19521568
..      ...
401     Huawei      P30 Lite      128    4         329       8038292
402     Samsung  Galaxy Note20 5G      128    8        1049      25629692
403     Xiaomi      Mi 10 Lite 5G      128    6         349       8526942
404     Apple    iPhone 12 Pro Max      128    6        1099      26851318
405     Oppo          Reno3      128    8         429      10481542
```

[347 rows x 6 columns]

**Feature Normalization** In this step, I use the zscore function to calculate the zscore of 'RAM' and 'Price (\$)'. These new values will help us to easily indicate whether the spec, price is below, near, or above the average. We decided to didn't calculate the zscore for Price in VND as it is just a calculated from other field and also has some rounding.

```
[ ]: result['RAM '] = zscore(result['RAM '])
      result['Price ($)'] = zscore(result['Price ($)'])
      result
```

```
[ ]:      Brand          Model  Storage  RAM  Price ($)  Price (VND)
0      Apple      iPhone 13 Pro      128  0.045702   1.909644      24408068
1      Samsung  Galaxy S21 Ultra      256  2.485475   2.558537      29294568
2      OnePlus          9 Pro      128  0.858959   1.585197      21964818
3      Xiaomi  Redmi Note 10 Pro      128  0.045702  -0.426371       6816668
4      Google          Pixel 6      128  0.858959   1.260751      19521568
..      ...
401     Huawei      P30 Lite      128 -0.767556  -0.264148       8038292
402     Samsung  Galaxy Note20 5G      128  0.858959   2.071867      25629692
```

403	Xiaomi	Mi 10 Lite 5G	128	0.045702	-0.199259	8526942
404	Apple	iPhone 12 Pro Max	128	0.045702	2.234090	26851318
405	Oppo	Reno3	128	0.858959	0.060298	10481542

[347 rows x 6 columns]

## 1.4 Data split

In this part, we decided to work with our diabetes dataset as the dataset is pretty much clean and we want to predict diabetes based on other attribute. Here we use all of the columns except the Outcome as x. And the Outcome column for Y. We split the data, 75% for training and 25% for testing

```
[ ]: diabetes = pd.read_csv('https://raw.githubusercontent.com/dinhphucv/CSC-177/
    ↪main/Data%20Preprocessing%20Project/data/diabetes.csv')
diabetes
```

```
[ ]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0              6        148             72           35         0  33.6
1              1         85             66           29         0  26.6
2              8        183             64            0         0  23.3
3              1         89             66           23        94  28.1
4              0        137             40           35       168  43.1
..          ...      ...             ...           ...         ...  ...
763            10        101             76           48       180  32.9
764             2        122             70           27         0  36.8
765             5        121             72           23       112  26.2
766             1        126             60            0         0  30.1
767             1         93             70           31         0  30.4
```

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[ ]: x = diabetes[diabetes.columns.difference(['Outcome'])]
y = diabetes['Outcome']
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25,
↳ random_state=12)
```

```
[ ]: print(x.columns)
```

```
Index(['Age', 'BMI', 'BloodPressure', 'DiabetesPedigreeFunction', 'Glucose',
      'Insulin', 'Pregnancies', 'SkinThickness'],
      dtype='object')
```

The below code is used to calculating the mean and std of certain columns in both training and testing set.

```
[ ]: feature_columns = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
↳ 'BMI', 'DiabetesPedigreeFunction']
for column in feature_columns:
    print(column)
    print('\tTraining set mean:', x_train[column].mean())
    print('\tTesting set mean:', x_test[column].mean())
    print('\tTraining set std:', x_train[column].std())
    print('\tTesting set std:', x_test[column].std())
```

Glucose

```
Training set mean: 120.52604166666667
Testing set mean: 122.0
Training set std: 31.30199630736171
Testing set std: 33.96625982181027
```

BloodPressure

```
Training set mean: 68.94965277777777
Testing set mean: 69.57291666666667
Training set std: 19.113006722286883
Testing set std: 20.110555584040394
```

SkinThickness

```
Training set mean: 20.73263888888889
Testing set mean: 19.947916666666668
Training set std: 15.694852459319465
Testing set std: 16.72905100325358
```

Insulin

```
Training set mean: 81.27256944444444
Testing set mean: 75.38020833333333
Training set std: 115.65762221613798
Testing set std: 114.17978027672025
```

BMI

```
Training set mean: 32.01597222222222
Testing set mean: 31.922395833333336
Training set std: 8.125176445501436
Testing set std: 7.131822974665879
```

DiabetesPedigreeFunction

```
Training set mean: 0.4719513888888889
Testing set mean: 0.47165104166666666
```

```
Training set std: 0.3358161055745912
Testing set std: 0.3183402606796537
```

#### 1.4.1 Mean and Std results (in case rerun change the data result):

Glucose

```
*****Training set mean: 120.52604166666667
*****Testing set mean: 122.0
*****Training set std: 31.30199630736171
*****Testing set std: 33.96625982181027
```

BloodPressure

```
*****Training set mean: 68.94965277777777
*****Testing set mean: 69.57291666666667
*****Training set std: 19.113006722286883
*****Testing set std: 20.110555584040394
```

SkinThickness

```
*****Training set mean: 20.73263888888889
*****Testing set mean: 19.947916666666668
*****Training set std: 15.694852459319465
*****Testing set std: 16.72905100325358
```

Insulin

```
*****Training set mean: 81.27256944444444
*****Testing set mean: 75.38020833333333
*****Training set std: 115.65762221613798
*****Testing set std: 114.17978027672025
```

BMI

```
*****Training set mean: 32.01597222222222
*****Testing set mean: 31.922395833333336
*****Training set std: 8.125176445501436
*****Testing set std: 7.131822974665879
```

DiabetesPedigreeFunction

```
*****Training set mean: 0.4719513888888889
*****Testing set mean: 0.47165104166666666
*****Training set std: 0.3358161055745912
*****Testing set std: 0.3183402606796537
```

### 1.5 Developing and documenting human insights with human interpretation on preprocessed data and possible effect on predictions.

We already learned that data quality is essential and poor data quality is an unfolding disaster. Here, we only deal with small datasets and already see why data preprocessing takes so much time. We must eliminate missing data, unnecessary features, duplicate data, outliers, or inconsistencies. After cleaning the data, our dataset is smaller, more manageable for us to read, and easier for the model to process. Data preprocessing ensures that we and the model are working with accurate and reliable information, which can lead to more accurate predictions.

## **1.6 Compare the two sets: the training data and the test data and analyze it, developing an intuition and meaning of your results.**

Splitting data is an essential step in machine learning. The training set will only be used to train, and the test set (unseen data) will be used to test how well a model is generalized. This process will prevent overfitting when the model performs well on the training set but poorly on the test set. In our case, the mean and std of columns from the two datasets are pretty close together.